

Julia for high-performance simulation of optical pulses

Corentin Simon

Université Libre de Bruxelles

2026-02-10

Simulations of optical pulse

In general, we want to integrate an equation of the form

$$\partial_z u(z, t) = \sum_n a_n \partial_t^n u + \widehat{N}(u)$$

⇒ Semilinear parabolic PDE

$$\partial_z u(z, t) = \hat{D}u + \widehat{N}(u)$$

- \hat{D} is a (stiff-linear) operator.
- \widehat{N} is some nonlinear function on u .

Stiffness problem

\hat{D} is stiff : **large eigenvalues**.

If transformed in the spectral domain $\partial_t \rightarrow -i\omega$

$$\hat{D}(\omega) = \sum_n a_n (-i\omega)^n$$

- Standard explicit ODE Method (RK45, ...) poorly suited due to stiffness.
- Implicit methods are better but required a nonlinear solver
 - Large system ($10^2 - 10^3$ or more points) \Rightarrow **poor performance !**

Split-step method

For the NLSE equation:

$$\partial_z u(z, t) = i \frac{\beta_2}{2} \partial_t^2 u + i\gamma |u|^2 u$$

Exact solution

- $u(L, t) = \exp(L(\hat{D} + \hat{N}))u(0, t)$
- Cannot expand the exponential as \hat{D} and \hat{N} do not commute.
- Uses the BCH approximation

$$\exp(h(\hat{D} + \hat{N}))u \approx \exp(h\hat{D}) \exp(h\hat{N}) + \mathcal{O}(h^2)$$

for small step h

- Split the step between linear and nonlinear one for a small
- Repeats L/h times.

Split-step implementation

For the NLSE,

- $\exp(h\hat{N})u_0 = \exp(i\gamma h |u_0|^2)u_0$

which can be directly computed

- For $\exp(h\hat{D})u_0$, \hat{D} is diagonal in the spectral representation

\Rightarrow exponentiation in the spectral domain

$$\Rightarrow \exp(h\hat{D})u_0 = \mathcal{F}^{-1}(\exp(h\hat{D}(\omega))\mathcal{F}(u_0))$$

Fourier transform for performing the DFT \mathcal{F} on u .

- Very performant $\mathcal{O}(n \log n)$
- Very fast/optimized implementation exists (MKL, FFTW, ...)

Split-step: strength and weakness

Strength

- Easy to implement

Weaknesses

- Lack accuracy $\mathcal{O}(h^2)$
- No adaptative control: fixed time step.
 - adaptative variations exist but are harder to implement.
- Required n exponential each step for the nonlinear part
 - Exponential is expensive
- For more complex nonlinear part \widehat{N} , no explicit formula for $\exp(\widehat{N})$
 - Numerical integration is required (RK45, ...)

RKIP - Stands for Runge-Kutta in the Interaction Picture

- Alternative method to split-step.
- Transform the stiff-problem into a non stiff one.
- Allow to use direct Runge-Kutta methods.

If we defined u_I as

$$u_I(z, t) = \exp(-t\hat{D})u$$

$$\begin{aligned}\partial_t u_I &= -\hat{D}u_I + \exp(-t\hat{D})\partial_t u \\ &= -\hat{D}u_I + \exp(-t\hat{D})(\hat{D}u + \hat{N}(u)) \\ &= -\hat{D}u_I + \hat{D}\exp(-t\hat{D})u + \exp(-t\hat{D})\hat{N}(u) \\ &= \exp(-t\hat{D})\hat{N}(u) = \exp(-t\hat{D})\hat{N}(\exp(t\hat{D})u_I)\end{aligned}$$

For a Runge-Kutta tableau $c_i, a_{ij}, \alpha_i, \tilde{\alpha}_i$, this gives

$$\begin{aligned} k_i &= \exp(-c_i h \hat{D}) \hat{N} \left(\exp(c_i h \hat{D}) \left(u_n + \sum_{j < i} h a_{ij} k_j \right) \right) \\ u_{n+1} &= \exp(h \hat{D}) \left(u_n + \sum_i h \alpha_i k_i \right) \text{ (Update)} \\ e &= \exp(h \hat{D}) \left(\sum_i h (\alpha_i - \tilde{\alpha}_i) k_i \right) \begin{array}{l} \text{(Error estimate} \\ \text{if adaptative)} \end{array} \end{aligned}$$

The best tableau for this method seems to be the Verner 5(6) tableau.

RKIP: strength and weakness

Strength

- As high order as the RK tableau
 - ▶ Large repertoire or higher-order efficient tables
- Adaptative with error control
- Agnostic about \hat{N}

Weaknesses

- Required computing $\exp(h\hat{D})$ each steps
 - ▶ If not adaptative, can be precomputed
 - ▶ If adaptative, can be cached for a fixed number of time steps h_1, h_2, \dots, h_S
→ restrict h to theses
- **Hard to implement**
 - ▶ Good news : **you don't have to !**

RKIP implementation

RKIP is now available as part of the Julia `OrdinaryDiffEq.jl` package.

Why Julia ?

- Julia is compiled Just-in-Time (JIT)
- Syntax close to Python and Matlab
- Faster interpreted language (even with tools/precompiled library)
- `OrdinaryDiffEq.jl` is the most performant ODE solver library available

How to use this code repo

1. Install Julia <https://julialang.org/downloads/> and git <https://git-scm.com/install/windows>

For Windows:

```
1 winget install --name Julia --id 9NJNWW8PVKMN -e -s msstore  
2 winget install --id Git.Git -e --source winget
```

shell

2. Ensure you have VSCode installed (<https://code.visualstudio.com/download>), with its Julia extension.
3. Clone the folder, and open it in VSCode.

```
1 git clone https://github.com/Azercoco/NLSE-Julia-Example.git  
2 cd NLSE-Julia-Example  
3 code .
```

shell

Once VSCode is open in the folder, open the command palette (Default : **Ctrl+Shift+P**)

Search for > **Julia** : Restart REPL, then execute with **Enter**

This will open the **Julia REPL** for interacting with Julia.

In the Julia REPL, type

```
1 ] # Open the Package manager  
2 activate .  
3 instantiate
```

julia

This will download and install all the required packages.

You can now play or use one of the examples.

Example 1: LLE

See `../examples/example_1_lle.jl`

This example shows how to solve the Lugiato-Leverefer equation

$$\partial_T u(T, t) = (-1 + i\partial_t^2)u + i(\Delta - |u|^2)u + S$$

which has two parameters S and Δ .

The operators are:

- $\hat{D} = -1 + i\partial_t^2$
- $\hat{N}(u) = i(\Delta - |u|^2)u + S$

Δ could have also been moved into \hat{D} but it is easier to include in \hat{N} as we want to be able to vary it with time $\Delta(t)$.

Dependency loading.

```
1 import ProgressLogging  
2 using Plots  
3 include("../src/mod.jl")
```

julia

- ProgressLogging is used to display a progress bar in VSCode and Plots for plotting.
- src/mod.jl contains all the routine used for the examples.

In Julia, a differential equation is defined as $f(u, p, t)$ where t is the integration variable and p a NamedTuple containing the equation parameters (here S and Δ).

```
1 p = (; S=1.1, Δ=1.1);
```

julia

This code create a NamedTuple with field S and Δ .

Crating the initial conditions.

```
1 u0 = lle_hss(p.S, p.Δ, 2^12; noise=1e-2)
```

julia

This creates a vector of size $2^{12} = 2048$ points, corresponding to the continuous stationary solution of the LLE equation with a random noise level of noise.

```
1 dτ = 0.05
2 LLE = SemilinearPDE(
3     StandartLLE(),
4     u0,
5     300.0,
6     dτ,
7     p
8 );
```

julia

This creates a semilinear PDE with initial solution u_0 , an integration bounds of $(0, 300.0)$, a time step $d\tau$ and parameters p .

We solve the problem

```
1 u = solve_rkip(LLE; abstol=1e-4, reltol=1e-8)
```

julia

with relative tolerance `reldtol` and absolute tolerance `abstol`. Changing these value will increase/decrease the precision with a cost/gain in performance.

We can plot the solution:

```
1 plot(  
2     LLE.τ,  
3     abs2.(u),  
4 )
```

julia

`LLE` has two fields `τ` and `freq` storing the value of fast time and normalized frequencies.

We can also sample the solution at different time using the `saveat` keyword arguments.

```
1 t = 0.0:5:300 # sample every 5 unit of time between 0 and 300
2 u2 = solve_rkip(LLE; saveat=t, abstol=1e-4, reltol=1e-8);
```

julia

`u2` is now matrix with a number of columns corresponding the size of `t`.

Similarly the solution is plotted.

```
1 heatmap(abs2.(u2))
```

julia

`LLE` has two fields `τ` and `freq` storing the value of fast time and normalized frequencies.

Exemple 2: LLE-Scan with time varying Δ

See :./examples/example_2_lle_scan.jl

This example is similar to the first. The only difference is that Δ is now a function of p and t .

```
1 Δ_scan(p, t) = p.Δ₀ + p.Δ_rate * t
2 p = ();
3 S=3.5,
4 Δ₀=-2.0,
5 Δ_rate=0.2,
6 Δ = Δ_scan,
7 )
8 u₀ = lle_hss(p.S, p.Δ₀, 1024; noise=1e-2)
```

julia

We have introduced two new parameters, the initial detuning Δ_0 and detuning scan rate Δ_rate . Otherwise, we solve it exactly the same manner as the first example.

Example 3: third-order dispersion

See :./examples/example_3_lle_d3.jl

In this example, we want to modify the linear operator \hat{D} to add a third-order dispersion to the LLE:

$$\partial_T u(T, t) = (-1 + i\partial_t^2 + id_3\partial_t^3)u + \dots$$

We add this new parameters d_3 to the set of parameters.

```
1 p = (;..., d3=0.08)
```

julia

```
1 LLE_d3 = SemilinearPDE(  
2     StandartLLE(;  
3          $\hat{D}=(\omega, p) \rightarrow -1 - 1im * \omega^2 + 1im * \omega^3 * p.d_3$   
4     ),...
```

julia

We can override the the definition of \hat{D} by passing a function of (ω, p) computing $\tilde{D}(\omega)$.

Example 4: Raman Nonlinearity

See `../examples/example_4_lle_rarman.jl`

The Raman nonlinearity adds an additional term to the LLE:

$$\partial_T u(T, t) = (-1 + i\partial_t^2)u + i(\Delta - |u|^2)u + S + i\tau_R(\partial_t|u|^2)u$$

Similarly, we add a new parameter `tau_R`.

```
1 p = (;..., τ_R=2e-4)
```

julia

And we can modify the nonlinearity used by doing

```
1 LLE_raman = SemilinearPDE(  
2     StandartLLE();  
3     ℑ=LLERamanNonlinearPart()  
4 ),...
```

julia

The rest is similar to the previous examples.

Adding a custom nonlinearity.

To add a new nonlinearity, declares a new empty struct corresponding this nonlinearity, which inherits AbstractNonlinearPart and implements (nl::CustomNonlinearity)(du, u, p, t).

```
1 struct CustomNonlinearity <: AbstractNonlinearPart end
2
3 function (nl::CustomNonlinearity)(du, u, p, t)
4     # Put your nonlinear function here and store the result in du
5 end
```

julia

and then use:

```
1 custom_lle = SemilinearPDE(
2     StandartLLE();
3     N=CustomNonlinearity()
4 ),...
```

julia

Custom Nonlinearity: PDLNSE

Example nonlinearity for a parametric driven NLSE:

$$\partial_T u(T, t) = (-1 + i\partial_t^2)u + i(\Delta - |u|^2)u + \kappa u^*$$

```
1 struct PDLNSE <: AbstractNonlinearPart end
2
3 function (nl::PDLNSE)(du, u, p, t)
4     Δ = get_var(nl, p.Δ, p, t)
5     κ = get_var(nl, p.κ, p, t)
6
7     @. du = lim * (abs2(u) - Δ) * u + κ * conj(u)
8 end
```

julia

The helper `get_var(nl, var, p, t)` allows `var` to be either a scalar or a function of `(p, t)`. We used the *broadcast* macro `@.` which executes the operation on the array without allocations.

Also remember to add `κ=...` to `p`.

Custom Nonlinearity: Cache

Sometimes we want to cache some array for the computation. In that case, we can implement function

`get_cache(nl, u0, p)` which returns a NamedTuple.

```
1 struct CustomNonlinearityWithCache <: AbstractNonlinearPart end
2
3 function get_cache(::CustomNonlinearityWithCache, u0, p)
4     return (; cached_array=similar(u0))
5 end
```

julia

The cache content can be accessed in the main evaluation with `p.cache`:

```
1 function (nl::PDLNSE)(du, u, p, t)
2     cached_array = p.cache.cached_array
3     ....
```

julia

Custom Nonlinearity: Raman example

Complete implementation in `../src/raman.jl`.

```
1 struct LLERamanNonlinearPart <: AbstractNonlinearPart end julia
2
3 # For Raman, we need to implement a cache to store intermediate computation
4 function get_cache(::LLERamanNonlinearPart, u0, p)
5     abs2_tmp = similar(u0) # pre-allocate an array for intermediate storage
6     raman_fourier_resp = convert(
7         typeof(u0),
8         @. -1im * 2π * p.freq * p.τ_R
9     ) # pre-computed the Rama kernel in spectral space
10    return (; abs2_tmp, raman_fourier_resp) # cache are supplied as a NamedTuple
11 end
```

```
1 @fastmath function (lle_raman::LLERamanNonlinearPart)(du, u, p, t) julia
2     S = get_var(lle_raman, p.S, p, t) # fetch eventually time varying function
3     Δ = get_var(lle_raman, p.Δ, p, t) # fetch eventually time varying function
4     @unpack abs2_tmp, raman_fourier_resp = p.cache # We can acces the cache using
5     p.cache
6     abs2_tmp .= abs2.(u) # used form temporary storage
7     @. du = 1im * (abs2_tmp - Δ) * u + S # LLE
8     fft!(p.bifft_plan, abs2_tmp) # in-place preallocated FFT
9     @. abs2_tmp = raman_fourier_resp * abs2_tmp # convolution = mutiplication in
10    spectral domain
11    ifft!(p.bifft_plan, abs2_tmp) # in-place preallocated IFFT
12    @. du += 1im * abs2_tmp * u
13    return du
14 end
```

To use `fft!` and `ifft` with pre-allocated storage and temporary array, the wrappers `ifft!/fft!` (`p.bifft_plan, target`) are available.

- Pre-allocate every array you may need in the cache (allocations are the main bottleneck of code)
 - Use Julia `similar(u0)` to create new array to keep to code generic.
 - Always use in-place mutating operation (in-place function in Julia are indicated with a !) which does not create temporary array
- Use the broadcast macro `@.` when to apply an element-wise operation on one or several arrays.
- Use the `@fastmath` macro before your function.
- Make sure all of your array have a well-defined type.
- Avoid using `abs(u)^2` to compute $|u|^2$, as it compute a squared-root internally. Use Julia `abs2()` instead
- Similarly, avoid `exp(1im*x)` to compute $\exp(ix)$ if x is real, use `cis(x)` instead.
- Be sure to computed only once and cached expensive computation that do not need to be updated.
- On IntelCPU, uses MKL FFT with `FFTW.set_provider!("mkl")`.
- Launch Julia with multiple threads wiht `julia --threads ...`
- For very large problem ($> 10^4$ points), consider using a GPU.
 - All the code in this repo should be compatible with a GPU by converting `u0` to `CuArray` from `CUDA.jl` package.
 - On the GPU, only use `Float32` (32-bit float) array instead of `Float64`.

Conclusion

- Except for simplicity, there is no reason to use Split-Step instead of RKIP
- Give Julia a try ! (You can interface it with Python)
- All examples code are on my GitHub <https://github.com/Azercoco/NLSE-Julia-Example>
 - ▶ Contributions and suggestions are welcome !